
Wording Rivals Peers in LLM Judicial Decision-Making

Austin L. Cheng
Stanford Law School
Stanford, CA 94305
austincheng98@gmail.com

Abstract

Language-model agents often converge after observing others’ decisions, but this pattern can reflect either peer influence or numerical anchoring. Convergence alone cannot distinguish them. We test this in fictional criminal sentencing with a yoked dose-response design. LLM judges see three numbers displaced from the guideline midpoint, attributed in paired conditions to peer sentences or a statistical forecast. Separate dose-response slopes estimate numerical pull under peer and forecast attribution, with their difference being the peer premium. With the guideline omitted, Claude Opus 5’s premium was 0.32 under bare blocks and 0.21 under structure-matched blocks. Under identical prompts and case allocations, GPT-6 Luna and Sol had matched premiums of 0.220 and 0.239 (both $p < 0.001$). Removing the guideline raises pull across all three Claude models. In civil damages, the estimated peer-minus-forecast difference in mean proportional distance from the displayed center was +0.140 for Opus 5 and +0.093 for GPT-6 Sol. Together, these results extend the peer premium to another model lineage and the peer-versus-forecast comparison to a separate legal domain, while showing that wording and guideline availability affect measured influence. The 4,086 core decision records, harness, analyses, and prompts are available at <https://github.com/austin-cheng98/llm-sentencing-agents>.

1 Introduction

Populations of language-model agents are increasingly used to study social processes alongside experiments with human participants. Simulated populations reproduce survey responses [4, 18], replicate classic behavioral experiments [1, 12], and model opinion dynamics on networks [8]. In these studies, agents who observe one another often move toward one another. Researchers often interpret this convergence as conformity, drawing on studies of group influence [5, 23] and information cascades [2, 6].

Convergence alone does not identify peer influence. When an agent sees a peer’s decision, its prompt contains both a social source and a number. Language models can move toward numbers presented in context regardless of their source [9, 14]; they can also move toward an interlocutor’s position [22]. Both mechanisms predict convergence. Prior studies distinguish them using discrete outcomes and varying whether a source is present [13, 27]. We instead use a continuous outcome and vary the displacement of displayed numbers. The resulting slope estimates numerical pull; the difference between source-specific slopes estimates the peer premium.

We study this distinction in fictional sentencing decisions. Paired conditions use the same cases and displayed numbers but attribute them either to other judges or to a statistical forecast. On Claude Opus 5 with the guideline removed, the premium is 0.32 under bare blocks and 0.21 under structure-matched blocks. Adding the routine instruction “You are deciding the same case independently” reduces it to -0.01 . The instruction changes pull by -0.28 ; its two paraphrases change pull by -0.33 and -0.20 . Weng et al. [26] raised independence rates with a persona instruction. Here, the routine sentence reduces the premium on Opus 5 but not on the two smaller models. Removing the guideline raises pull across the three Claude models. In a sequential cascade, 71% of decisions exactly repeat a predecessor, compared with a 48% exact-match rate when independent agents are shuffled into speaking orders.

Sentencing provides a continuous outcome, an institutional anchor that can be removed, and prior evidence that human sentencing decisions respond to numerical anchors [10, 11]. Language models have also been proposed as drafting aids in this setting.

2 Related work

Conditioning a model on demographic profiles reproduces human subgroup response distributions [4, 18]. Agent architectures with memory and reflection produce social interactions [17], and classic experiments have been replicated with simulated subjects [1, 12]. However, Santurkar et al. [20] document systematic gaps between model opinion distributions and their targets. Model outputs also change with wording: Sclar et al. [21] show that formatting alone can change benchmark accuracy, and Mizrahi et al. [16] argue that single-template evaluation misstates model rankings. We ask whether a routine instruction changes the measured size of a peer effect, rather than task accuracy. Chuang et al. [8] report that LLM agents in network simulations converge toward accurate consensus; prompts inducing confirmation bias produce fragmentation. Zhou et al. [28] find that collective effects can vanish or reverse once validity conditions are enforced. Two non-social mechanisms could produce convergence: numeric anchoring [9, 14] and movement toward an interlocutor’s expressed views [22]. In humans, Englich et al. [10] found experienced judges shifting sentencing recommendations toward dice-roll anchors, and Guthrie et al. [11] documented anchoring among United States federal magistrate judges.

3 Experimental design

3.1 Task and cases

Decisions come from Claude Opus 5, Sonnet 5 or Haiku 4.5 at default decoding settings, one draw per design cell, and a GPT-6 Luna sample (Section 4.9). Repeat draws on a byte-identical prompt, available for 56 cells collected twice, differ by a within-cell standard deviation of 0.064.

Each agent acts as a sentencing judge in a fictional jurisdiction. It returns an integer sentence in months, a confidence rating, a primary and a secondary rationale from a fixed six-item menu, and a short justification. Every prompt states that the cases are procedurally generated and that no real defendant or court is involved.

We generate 16 cases by fully crossing four binary factors: offense severity, prior record, expressed remorse, and substantial assistance to investigators. Offense type rotates across four families, so each type co-occurs equally often with each factor level. Severity and prior record determine an advisory range taken from the United States Sentencing Guidelines table, giving four ranges below the statutory maxima. Remorse and assistance do not appear in the range, so they can affect sentences only through departures from the guideline.

Before deciding, agents in the anchor arms see three numbers centred on a displaced point,

$$\bar{a}_{ic} = m_c(1 + \delta_{ic}), \quad \delta_{ic} = D[(g_c + i) \bmod 4], \quad D = (-0.30, -0.15, +0.15, +0.30), \quad (1)$$

where m_c is the midpoint of case c 's advisory range, $g_c \in \{0, 1, 2, 3\}$ is its design group and i indexes the agent.

Appendix D gives the full rotation in Equation 1. The sixteen design points are partitioned into four groups in which every case factor splits evenly. Each agent therefore sees every displacement exactly four times, and the four factors split evenly within each displacement level. As a result, $\hat{\pi}$ carries no part of a factor effect. The rotation has period four, so with six agents the first two displacement columns repeat.

3.2 Conditions and estimation

We cross attribution with the presence of an institutional anchor. In the peer condition, the numbers are sentences already entered in this case by other judges of the same bench. In the forecast condition, they are the output of a regression model fitted to historical dispositions. In the docketing condition, they are stored values attached to the file by the case-management system, with no legal significance. The guideline-present condition states the advisory range; the guideline-removed condition states only the statutory maximum. Each decision uses a fresh context, preventing information carryover across prompts.

Each block originally ended with a closing sentence, but the sentences were not matched. The peer block closed with "You are deciding the same case independently," whereas the forecast block ended with a note that the forecast is advisory. We added arms with both closing sentences removed, an arm that gave the peer block a descriptive sentence matching the forecast block's structure, and two arms with paraphrases of the independence sentence. The artifact repository's `experiment/prompts.txt` file gives all nine prompts verbatim.

The outcome is the proportional deviation of the sentence from the case's guideline midpoint, $y_{ijc} = (s_{ijc} - m_c)/m_c$, which puts cases with different ranges on a common scale. The quantity of interest is the pull π ,

$$y_{ijc} = \alpha + \pi \delta_{ic} + \gamma^\top F_c + \varepsilon_{ijc}, \quad (2)$$

where F_c collects the four case factors. $\pi = 0$ means the agent ignores the numbers and $\pi = 1$ means it adopts them in full. Writing $P_j = 1$ for the peer arm and 0 for the forecast arm, the peer premium Δ is the difference in slopes recovered by interacting P_j with δ_{ic} in Equation 2,

$$y_{ijc} = \alpha + \pi \delta_{ic} + \beta P_j + \Delta (\delta_{ic} \times P_j) + \gamma^\top F_c + \varepsilon_{ijc}, \quad (3)$$

so that $\Delta = \partial^2 \mathbb{E}[y] / \partial \delta \partial P$. A moderation model adding a guideline indicator is reported in Appendix E; only the arms carrying the original closing lines were run on both sides of it, so we report the two guideline conditions separately instead.

Paired arms share agents, cases, and displacements, so each cell (i, c) contributes one decision under each label. Under the sharp null that a decision would have been the same under either label, $y_{ic}(P=1) = y_{ic}(P=0)$ for every cell, the labels within a cell are exchangeable; relabelling them at random per cell generates the exact null distribution of $\hat{\Delta}$, from which we calculate a two-sided p -value and a 95th-percentile critical value over $B = 9,999$ draws. Each contrast has its own critical value because it rests on a different set of cells. The values range from 0.12 to 0.18 here, and each premium is evaluated against its corresponding critical value. Because the contrast never crosses agents, between-agent heterogeneity cancels from the null rather than inflating it. Standard errors cluster on case, agent, and displacement cell. Appendix E gives the relabelling scheme, the wild cluster bootstrap used for cluster-robust p -values, the per-arm displacement null, the between-agent spread $\bar{\sigma}$, and the yoke, factor-balance, and offense-type checks.

3.3 Civil-damages extension

The extension tests monetary valuations on 16 fictional bodily-injury claims in a 2^4 factorial crossing severity, liability clarity, claimant mitigation, and documentation completeness. Each claim has a stipulated midpoint and an advisory range of $\pm 12.5\%$, but the range is withheld; agents may enter

any amount up to the one-million-dollar policy limit. Prompts show three valuations at 94%, 100%, and 106% of a midpoint shifted by $\delta \in \{-0.30, -0.15, 0.15, 0.30\}$, rounded to the nearest hundred dollars. Matched primary arms use identical cases and figures and the same no-other-information cue; only attribution (other adjusters or a regression forecast) changes. We measure in-band awards, proportional distance from the displayed center, awards within 2% of center, and exact matches. D1’s initial Opus 5 panel informed outcome selection; the measures and directions were registered before examining two additional agents (32 pairs). Preregistered D2 reused the same 320 prompts and assignments with GPT-6 Sol high; one failed response left 159 pairs.

The in-band indicator includes both displayed endpoints. Proportional distance is the absolute deviation from the central figure divided by that figure; we also code whether the award falls within 2% of center and whether it matches the center exactly. Binary outcomes use two-sided exact sign tests on discordant matched pairs. For distance, we estimate the mean paired peer-minus-tool difference, cluster uncertainty by claim, and report the registered minimum detectable effect.

4 Results

4.1 The peer premium

On Claude Opus 5 with bare blocks, peer attribution gives a pull of 0.88, compared with 0.56 for the forecast, yielding a premium of 0.32 (SE 0.05 on case, 0.09 on agent, $p < 0.001$, $n = 128$). The peer block has four lines, whereas the forecast block has five because naming a forecast requires a descriptive sentence. Matching the structures reduces the premium to 0.21 (SE 0.08 on case, $p = 0.005$), against a design null of 0.15; restoring the closing sentences gives -0.01 . The lower structure-matched estimate is consistent with the unmatched block structures contributing to the bare comparison. The bare comparison on two smaller models gives 0.12 and 0.02; the last four rows report the GPT-6 cross-lineage estimates (Section 4.9; Figure 1).




Figure 1: Every estimate compares peer attribution against a statistical forecast over bit-identical numbers, with the advisory guideline removed throughout. Shading marks the design null of that row’s own contrast, the 95th percentile of its label-swap distribution. Band widths vary because the contrasts have different design-specific null distributions. GPT-6 rows have no label-swap band. Intervals are 95% and cluster on case.

Across bare-block arms, premiums are 0.32 on Opus 5, 0.12 on Haiku 4.5, and 0.02 on Sonnet 5. The ordering reverses with the guideline: Opus 5 is at -0.02 , versus 0.20 and 0.21 in cells containing 32 decisions from one agent (Appendix G). These cells are too small to rank the models. The independent-decision sentence reduces the Opus 5 premium but not those in the smaller-model arms, so comparisons depend on wording.

4.2 Instruction wording

“You are deciding the same case independently” changes pull by -0.28 ($p = 0.002$, $n = 128$); “Form your own view of what this case warrants” changes it by -0.33 ($p < 0.001$, $n = 96$); and

“Your sentence should reflect your own judgment of the case” changes it by -0.20 ($p = 0.016$). Each change is statistically distinguishable from its corresponding label-swap null. None names a peer, withdraws the block, or changes a displayed number. The forecast hedge leaves pull unchanged (0.05 , $p = 0.400$). Because a premium is a difference in pulls, these effects account for the premium observed with both original sentences in place:

$$\hat{\Delta}_{\text{orig}} = \hat{\Delta}_{\text{bare}} + \hat{c}_{\text{peer}} - \hat{c}_{\text{fcst}} = 0.32 + (-0.28) - (0.05) = -0.01, \quad (4)$$

Because the paired prompts keep cases and displayed numbers fixed, this decomposition estimates the response to wording changes within this task. It does not distinguish whether the added sentence changes perceived independence, the social meaning assigned to peer sentences, or attention to the anchors. The near-zero forecast-hedge estimate provides a comparison: adding a sentence did not reliably change pull in that arm. The results locate a wording effect in the peer condition while leaving its psychological mechanism unmeasured (Figure 2).




Figure 2: Pull before and after the sentence is appended, with the numbers shown to the agent held identical throughout.

Measured against the same bare forecast, the two paraphrases have different effects: one reduces the estimated premium to 0.05, whereas the other leaves it at 0.18 (Appendix H). All three wordings lower peer pull, which points to the instruction’s content rather than its exact wording.

The sentence also changes justifications: peer mentions fall from 42% to 6% when it is added, alongside the reduction in pull, although the text does not establish how agents interpreted the instruction.

4.3 The advisory guideline

Removing the advisory range raises estimated pull in each of the three Claude models. The pooled interaction is 0.54 (SE 0.07, $p < 0.001$). By model, the estimates are 0.54 for Opus 5, 0.63 for Sonnet 5, and 0.36 for Haiku 4.5 (Figure 3, Appendix G). Only arms carrying the original closing lines were run with the guideline, so this contrast holds those lines fixed rather than spanning the wordings in Section 4.2. Pull under a guideline is 0.07, 0.24, and 0.56. Agent fixed effects over the 6 agents in both settings leave the interaction at 0.54 (SE 0.07).




Figure 3: Pull with 95% intervals clustered on case. (a) Original closing sentences, with (solid) and without (hatched) the advisory guideline. (b) Bare blocks, no guideline, collected for every model.

Between-agent spread is 0.029 with the guideline and 0.286 without it. Leave-one-out estimates for the no-guideline condition range from 0.074 to 0.327 (Appendix G). Without a guideline, mean deviation is 0.26 compared with -0.20 ; displaying the three numbers shifts it to -0.03 .

4.4 The docketing arm

A docketing condition provides a source comparison with no legal significance. Presented as case-management records, the same values move agents by 0.07 (SE 0.09, $n = 64$), compared with 0.61 for the forecast (-0.54 , SE 0.11, $p < 0.001$). The prompt also tells agents the values are unrelated to the merits, so the estimate combines the source label with the instruction to disregard the numbers (Figure 4).




Figure 4: Small points are decisions; large points are cell means; the diagonal marks full adoption. (a) Opus 5, original blocks, guideline present (solid) or removed (hatched); (b) Opus 5, bare blocks; (c) GPT-6 Luna and Sol pooled within each bare-block arm, with model-specific intercepts and a common slope; (d) selected pull estimates with 95% intervals clustered on case.

Case-factor estimates show that sentence length varies with case characteristics even when no numbers are shown (Appendix L). Severity and prior record have larger effects than assistance and remorse. With a guideline, most sentences fall within its range; without it, proportionality is less commonly named and deterrence is more common.

4.5 Confidence and justifications

Each decision has a confidence rating and a written justification. Mean confidence is negatively correlated with pull across the six unguided anchor arms of Claude Opus 5 ($r = -0.82$; Figure 8, Appendix H). The bare peer arm has the highest pull and lowest confidence, at 6.00; the docket-

ing arm has the lowest pull and highest confidence, at 6.84. Claude Sonnet 5 and Haiku 4.5, and GPT-6 Luna and Sol, are plotted separately for their available shared arms; only Opus 5 enters the fit. Panel 8b averages exact adoption by provider across the three arms shared by all five models. The correlation is descriptive and based on six arm means; confidence does not identify decisions unaffected by the manipulation.

Written justifications differ by instruction even when estimated pull is similar. With the guideline removed and each block carrying its own closing sentence, peer and forecast pull are 0.61 and 0.61 (difference -0.01 , $p = 0.945$). The peer arm names displayed numbers in 6% of justifications, compared with 64% in the forecast arm. Without the peer closing line, peer mentions occur in 42% of justifications; with the line, they occur in 6%. These differences are consistent with the instruction affecting what agents report, while the pull estimates remain similar.

Across decisions, pooling every unguided anchor arm, justifications that name the displayed numbers accompany a pull of 0.45 compared with 0.67 for justifications that do not. This comparison is descriptive because it separates decisions rather than randomized conditions. This association echoes Turpin et al. [24]’s finding that model explanations may not reveal what influenced an output, though that study examined single-agent reasoning.

4.6 The cascade

Earlier arms used experimenter-selected numbers. The cascade measures copying under sequential observation: six agents decide the same case one after another, and each sees the sentences produced by earlier agents. Among the 80 decisions with a visible predecessor, 71% reproduce one predecessor’s number exactly (binomial interval: 61 to 81%). In 5 of sixteen cases, the cascade ends unanimously. The comparison floor is 48%, obtained by shuffling independent agents into pseudo-speaking orders in which influence is impossible. Between-agent spread falls from 0.056 across the first two speakers to 0.017 across the final pair. Mean deviation is 0.19 at the first speaking position and 0.12 at the last. The first speaker receives no peer block and matches the no-peer arm within 0.06 (SE 0.04), so the cascade compares conditions within one arm (Figure 5).




Figure 5: (a) Each line is one case, centred on its own mean, with agents plotted by speaking position; the heavy line is mean absolute deviation. (b) Between-agent spread on the same cases, for independent agents and for the first and last pairs of a cascade.

Both cascade statistics have a comparison floor, which determines the primary estimator. With case fixed effects, the observed slope is 0.55 (SE 0.06), compared with a placebo floor of -0.97 . The plain slope is close to its own floor, so we use the fixed-effect version as primary. Appendix I gives both floors and the per-seat copying rates.

4.7 Round numbers

Sentence lengths cluster at round values. Over 3807 decisions agents use 93 distinct values; 65% are divisible by six months and 42% by twelve. A decision can therefore match a displayed number because of this clustering rather than copying. The exact-match rate is 40% when at least one

displayed number is round, compared with 18% when none is. We therefore compare copying with a shuffled floor that preserves this preference. The peer premium is unaffected because both arms see identical numbers. Only 16.7% of the displayed numbers are divisible by six, approximately the chance rate.

Among records in which no displayed number is divisible by six, the bare premium is 0.35 ($p = 0.007$, $n = 64$), compared with 0.32 in the full sample. The structure-matched premium is 0.19 ($p = 0.238$, $n = 96$), compared with 0.21. The point estimates are similar, but the matched estimate is not statistically significant in this half-sample. The log-ratio and within-case-rank estimates remain positive for both the bare and structure-matched premiums (Appendix J). Confidence uses only 4 of its ten values, with 98% of ratings at six or seven; Section 4.5 reports arm-mean patterns, not a calibrated measure.

4.8 Equivalence and reliability

The structure-matched premium could reflect attribution or perceived informativeness: judges may view peer sentences as drawing on more case information than historical forecasts. A preregistered pair gave both blocks the same statement: “These figures were produced from the same case file you have been given, and from no other information about this defendant or this offense.” Headers and row labels were the only remaining contrasts; prompts and analyses were fixed before the 128 decisions.

The premium is -0.005 ($p = 0.973$, $n = 128$; 95% CI $[-0.247, 0.238]$). At 80% power and a two-sided 5% test, the MDE is 2.80 times the clustered SE. Here it is 0.347, above the 0.21 benchmark, and the interval contains that value. The registered stopping rule therefore leaves the estimate imprecise; the arm was not extended after outcome inspection.

The cue also changes number use: exact reproduction falls by roughly two thirds, out-of-range sentences rise to 80–84%, and dispersion increases (Appendix K). Agents say the figures add no case information because they derive from the same file. The result combines source attribution with the no-other-information instruction and does not identify the source effect when agents use the numbers more.

A separate GPT-6 Luna follow-up used neutral access wording, describing both sources as having reviewed the case file without saying they used no other information. Before inspecting outcomes, we fixed a factor-balanced sample of 12 cases across 26 contexts (312 matched pairs). The premium is $+0.250$ ($p < 0.001$, 95% CI $[+0.134, +0.365]$; MDE 0.161); its interval excludes zero but contains the 0.206 benchmark. Number use was higher under peer attribution: slopes 0.911 vs. 0.661, exact matches 50.6% vs. 15.4%, and out-of-range sentences 23.1% vs. 71.8%. This distinct prompt comparison does not resolve the registered no-other-information contrast.

A second registered pair added the same reliability claim to both blocks. Its premium is 0.001 ($p = 0.995$, $n = 128$, MDE 0.390); adding the claim changes peer-arm pull by 0.012 ($p = 0.905$) and forecast-arm pull by 0.006 ($p = 0.945$). All four estimates are near zero and underpowered against 0.21; they do not establish whether the claim restores the premium.

4.9 Other models

The structure-matched premiums on Sonnet 5 and Haiku 4.5 are imprecise against the Opus benchmark (-0.142 , MDE 0.622; -0.012 , MDE 0.219). Haiku’s bare premium is 0.121 ($p = 0.036$); these smaller-model cells do not support reliable rankings.

GPT-6 Luna and Sol each contribute 1,248 decisions across 26 contexts, using identical prompts, cases, numbers, and assignments (Table 2). Their matched premiums are 0.220 and 0.239; bare premiums are 0.256 and 0.113 (all $p < 0.001$). The MDEs are below the corresponding Opus benchmarks. The matched intervals include those benchmarks, while Sol’s bare interval is below the Opus estimate. Luna used extra-high and Sol high reasoning; their estimate difference is descriptive and does not isolate model identity.

4.10 Civil-damages extension

In the registered Opus 5 confirmation, 14/32 peer awards fell within the range shown in the prompt, compared with 26/32 forecast awards (exact sign test, $p = 0.0042$). Peer awards were also less likely to fall within 2% of the center (10/32 vs. 24/32; $p = 0.00052$).

Exact matches were less frequent under peer attribution (2/32 vs. 23/32; $p < 0.000001$). Mean proportional distance from the center was 0.217 for peer awards and 0.077 for forecast awards (difference +0.140, 95% CI [0.032, 0.248]), just below the registered MDE of 0.142.

The full D1 panel is descriptive because its first eight agents informed outcome selection. Across that panel, peer awards were less close to the center on all three measures; mean distance was 0.392, compared with 0.156 for forecast awards (difference +0.236, 95% CI [0.145, 0.328]).

Preregistered D2 used the same 320 prompts and assignments with GPT-6 Sol high. One failed response left 159 matched pairs. In-band rates were similar (90/159 vs. 88/159; $p = 0.871$), as were rates within 2% of center and exact matches (Figure 6; $p = 1.000$ and $p = 0.200$). Mean distance remained higher under peer attribution: 0.274 vs. 0.182 (difference +0.093, 95% CI [0.011, 0.174]). This estimate was positive but below the registered MDE of 0.107. Thus, D2 did not reproduce the in-band contrast, although its mean-distance estimate pointed in the same direction as Opus 5 (Figure 6). Appendix N reports full results.




Figure 6: Civil-damages results. Top: discordant matched pairs as a share of all pairs, with peer-only awards extending right (blue) and forecast-only left (red). Bottom: peer-minus-forecast distance from center with 95% claim-clustered t intervals; open diamonds mark minimum detectable effects. Samples: Opus 5 D1 confirmation ($n = 32$), full D1 panel ($n = 160$, descriptive), and GPT-6 Sol D2 ($n = 159$).

5 Discussion

The central estimand is the response to a displayed number when its source attribution changes, holding the case and number fixed. For Opus 5, matching block structure reduces the bare-block premium from 0.32 to 0.21. The matched comparison removes differences in length and closing sentences, making attribution a more direct interpretation of the contrast. It still measures a response to source labels in this prompt, rather than identifying conformity as the agent's psychological mechanism. The lower matched estimate also shows why the wording of both alternatives is part of the estimand.

Other prompt changes locate further sources of sensitivity. The independent-judgment sentence and its two paraphrases reduce peer pull, while the forecast hedge leaves pull nearly unchanged. Remov-

ing the advisory guideline raises pull across the three Claude models (Figure 3b), indicating that the institutional reference point changes how agents use the displayed values. Docketing labels produce little pull when agents are instructed to disregard those values. In the cascade, exact reproduction (71%) exceeds the shuffled baseline (48%), while between-agent spread falls across speaking positions. Together, these results distinguish responses to source, wording, institutional anchors, and sequential exposure; no single convergence measure captures all four. GPT-6 Luna and Sol also show positive premiums, extending the comparison to another model lineage without establishing equivalence across models.

The civil-damages extension tests the same source contrast on fictional compensation awards. In the registered Opus 5 confirmation, 14/32 peer awards, compared with 26/32 forecast awards, fell within the displayed range. GPT-6 Sol rates were similar (90/159 vs. 88/159), so the in-range contrast did not replicate. Mean proportional distance from the displayed center was higher under peer attribution in both runs (+0.140 for Opus 5; +0.093 for GPT-6 Sol). The Sol estimate was positive (95% CI [0.011, 0.174]) but below the registered MDE (0.107). Thus, the in-range difference appears only in Opus 5, while mean distance is higher under peer attribution in both runs. This extension applies the design to a second legal outcome, but does not show a uniform pattern across models and measures.

For legal simulations, these findings make source wording and institutional anchors measurable features of the evaluation. A simulated peer effect should be reported separately from accuracy on held-out decisions. The same distinction applies when simulated agents are used to study social influence. Matching the displayed information and textual structure, then varying source and context, helps identify which prompt component changes the measured response.

6 Limitations

The sentencing cases and civil claims are fictional, and the study has neither a human baseline nor an external measure of decision accuracy. Its primary premium estimates how outputs respond to attribution under controlled prompts. They do not measure whether agents use legally relevant facts. The structure-matched estimate is also specific to the wording and information statement tested here. Other source descriptions or institutional settings could produce different responses. The no-other-information wording changes number use as well as attribution, and the neutral-access follow-up uses different wording; neither comparison isolates a general source effect across all information conditions.

Precision varies across arms. Equivalence and reliability comparisons remain underpowered, and smaller-model estimates are too imprecise for dependable rankings. The Claude models share a lineage; GPT-6 is a separate family, with Luna and Sol evaluated under different reasoning settings. Each prompt cell uses one draw, so the estimates do not characterize run-to-run variability. In the civil-damages study, the full Opus 5 D1 panel (160 pairs) is descriptive because its first eight agents informed outcome selection; the registered confirmation uses 32 pairs. GPT-6 Sol reused the 320 prompts and assignments, with one failed response leaving 159 pairs. Its in-range rates were similar (90/159 peer vs. 88/159 forecast), but its mean proportional-distance difference was positive (+0.093, 95% CI [0.011, 0.174]) and below the registered MDE (0.107). The interval excludes zero, although the point estimate is smaller than the preregistered detection threshold. The cascade also uses a single six-agent sequence per case, which limits inference about other group sizes, speaking orders, and exposure patterns. These design and precision limits bound the claims to the studied prompts, models, and fictional cases.

7 Conclusion

Across fictional sentencing decisions, measured peer premiums vary with source attribution, prompt wording, and guideline availability. Structure matching clarifies the source comparison, and positive GPT-6 estimates extend it to another model lineage. The civil-damages extension applies the design to fictional compensation awards. The in-range difference appears in the Opus 5 confirmation but

not GPT-6 Sol; mean distance is higher under peer attribution in both. These results do not establish a uniform effect across measures and models. The experiments measure prompt sensitivity. Evaluations of legal AI should distinguish responses to recommendation sources from accuracy against case outcomes.

References

- [1] Gati V. Aher, Rosa I. Arriaga, and Adam Tauman Kalai. Using Large Language Models to Simulate Multiple Humans and Replicate Human Subject Studies. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, volume 202, pages 337–371, 2023.
- [2] Lisa R. Anderson and Charles A. Holt. Information Cascades in the Laboratory. *American Economic Review*, 87(5):847–862, 1997.
- [3] Joshua D. Angrist. The Perils of Peer Effects. *Labour Economics*, 30:98–108, 2014. doi: 10.1016/j.labeco.2014.05.008.
- [4] Lisa P. Argyle, Ethan C. Busby, Nancy Fulda, Joshua R. Gubler, Christopher Rytting, and David Wingate. Out of One, Many: Using Language Models to Simulate Human Samples. *Political Analysis*, 31(3):337–351, 2023. doi: 10.1017/pan.2023.2.
- [5] Solomon E. Asch. Studies of Independence and Conformity: I. A Minority of One Against a Unanimous Majority. *Psychological Monographs*, 70(9):1–70, 1956. doi: 10.1037/h0093718.
- [6] Sushil Bikhchandani, David Hirshleifer, and Ivo Welch. A Theory of Fads, Fashion, Custom, and Cultural Change as Informational Cascades. *Journal of Political Economy*, 100(5):992–1026, 1992. doi: 10.1086/261849.
- [7] A. Colin Cameron, Jonah B. Gelbach, and Douglas L. Miller. Bootstrap-Based Improvements for Inference with Clustered Errors. *Review of Economics and Statistics*, 90(3):414–427, 2008. doi: 10.1162/rest.90.3.414.
- [8] Yun-Shiuan Chuang, Agam Goyal, Nikunj Harlalka, Siddharth Suresh, Robert Hawkins, Sijia Yang, Dhavan Shah, Junjie Hu, and Timothy T. Rogers. Simulating Opinion Dynamics with Networks of LLM-based Agents. In *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 3326–3346, 2024. doi: 10.18653/v1/2024.findings-naacl.211.
- [9] Jessica Maria Echterhoff, Yao Liu, Abeer Alessa, Julian McAuley, and Zexue He. Cognitive Bias in Decision-Making with LLMs. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 12640–12653, 2024. doi: 10.18653/v1/2024.findings-emnlp.739.
- [10] Birte Englich, Thomas Mussweiler, and Fritz Strack. Playing Dice with Criminal Sentences: The Influence of Irrelevant Anchors on Experts’ Judicial Decision Making. *Personality and Social Psychology Bulletin*, 32(2):188–200, 2006. doi: 10.1177/0146167205282152.
- [11] Chris Guthrie, Jeffrey J. Rachlinski, and Andrew J. Wistrich. Inside the Judicial Mind. *Cornell Law Review*, 86(4):777–830, 2001.
- [12] John J. Horton. Large Language Models as Simulated Economic Agents: What Can We Learn from Homo Silicus? *NBER Working Paper 31122*, 2023. doi: 10.3386/w31122.
- [13] Yibo Hu and Jiaming Qu. Most LLM Conformity Needs No Speaker: Measuring the Speaker-Free Floor in Peer-Pressure Benchmarks. *arXiv preprint arXiv:2607.05545*, 2026. doi: 10.48550/arXiv.2607.05545.
- [14] Erik Jones and Jacob Steinhardt. Capturing Failures of Large Language Models via Human Cognitive Biases. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, pages 11785–11799, 2022.

- [15] Charles F. Manski. Identification of Endogenous Social Effects: The Reflection Problem. *The Review of Economic Studies*, 60(3):531–542, 1993. doi: 10.2307/2298123.
- [16] Moran Mizrahi, Guy Kaplan, Dan Malkin, Rotem Dror, Dafna Shahaf, and Gabriel Stanovsky. State of What Art? A Call for Multi-Prompt LLM Evaluation. *Transactions of the Association for Computational Linguistics*, 12:933–949, 2024. doi: 10.1162/tacl_a_00681.
- [17] Joon Sung Park, Joseph C. O’Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative Agents: Interactive Simulacra of Human Behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology (UIST)*, 2023. doi: 10.1145/3586183.3606763.
- [18] Joon Sung Park, Carolyn Q. Zou, Aaron Shaw, Benjamin Mako Hill, Carrie Cai, Meredith Ringel Morris, Robb Willer, Percy Liang, and Michael S. Bernstein. Generative Agent Simulations of 1,000 People. *arXiv preprint arXiv:2411.10109*, 2024. doi: 10.48550/arXiv.2411.10109.
- [19] David Roodman, Morten Ørregaard Nielsen, James G. MacKinnon, and Matthew D. Webb. Fast and Wild: Bootstrap Inference in Stata Using boottest. *The Stata Journal*, 19(1):4–60, 2019. doi: 10.1177/1536867x19830877.
- [20] Shibani Santurkar, Esin Durmus, Faisal Ladhak, Cino Lee, Percy Liang, and Tatsunori Hashimoto. Whose Opinions Do Language Models Reflect? In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, volume 202, pages 29971–30004, 2023.
- [21] Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. Quantifying Language Models’ Sensitivity to Spurious Features in Prompt Design, or: How I Learned to Start Worrying about Prompt Formatting. In *The Twelfth International Conference on Learning Representations (ICLR)*, 2024.
- [22] Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. Towards Understanding Sycophancy in Language Models. In *The Twelfth International Conference on Learning Representations (ICLR)*, 2024.
- [23] Muzafer Sherif. A Study of Some Social Factors in Perception. *Archives of Psychology*, 27 (187):1–60, 1935.
- [24] Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language Models Don’t Always Say What They Think: Unfaithful Explanations in Chain-of-Thought Prompting. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [25] Matthew D. Webb. Reworking Wild Bootstrap-Based Inference for Clustered Errors. *Canadian Journal of Economics*, 56(3):839–858, 2023. doi: 10.1111/caje.12661.
- [26] Zhiyuan Weng, Guikun Chen, and Wenguan Wang. Do as We Do, Not as You Think: the Conformity of Large Language Models. *International Conference on Learning Representations (ICLR)*, 2025.
- [27] Huixin Zhong et al. Disentangling the Drivers of LLM Social Conformity: An Uncertainty-Moderated Dual-Process Mechanism. *arXiv preprint arXiv:2508.14918*, 2025. doi: 10.48550/arXiv.2508.14918.
- [28] Jiaxu Zhou et al. The PIMMUR Principles: Ensuring Validity in Collective Behavior of LLM Societies. *arXiv preprint arXiv:2509.18052*, 2025. doi: 10.48550/arXiv.2509.18052.

A Ethical considerations

We simulate criminal sentencing, a setting where decisions affect defendants’ liberty. Every case is procedurally generated and fictional, and every prompt says so. This study does not evaluate language models as sentencing tools. With the guideline removed, model output responds to the arbitrary forecast displacement (0.61). The sentences are measurements of model behavior on synthetic cases, not judgments about real defendants.

B The memory experiment

Deployed agent architectures more often show an agent a summary of itself than a summary of its neighbours, so we ran a second experiment on the same cases. Agents in the own-record arm receive a profile of their past decisions, refreshed every four cases, giving their mean sentence, their average departure from the advisory range, how their sentences differed by remorse and by assistance, and their most frequent rationale. Agents in the no-memory arm decide the identical cases in the identical order with a fresh context each time, and four baseline cases return at the end in paraphrased form, which gives a within-agent test that no change in case composition can produce.

Without memory, the estimated trend across sixteen cases is 0.01 (SE 0.06). With an own-record summary, the estimated trend is 0.09 (SE 0.09), which is not distinguishable from zero. The paired difference between the arms on the same case at the same position is 0.006 (SE 0.011). On returning cases, the estimated shifts are -0.01 without memory and 0.01 with memory.

We tested own-record summaries because they provide numerical information similar to peer-supplied sentences. Each arm contains three agents and sixty decisions, limiting precision for smaller effects.

C What was collected

The original arms were drawn on 17–18 August 2026 at default decoding settings, one draw per cell, using only the prompts in the artifact repository’s `experiment/prompts.txt` file; staged collection and later agent additions explain the unequal arm sizes.

The registered cross-model arms were collected from 30 September to 1 October through the released harness, which creates a fresh context for each prompt and inherits decoding settings. Parse failures were recorded without redraws; one Haiku structure-matched cell failed, leaving 63 rather than 64 decisions. The GPT-6 Luna sample used 26 fresh contexts across three arms and 16 cases (1,248 decisions). It used isolated contexts with model `gpt-6-luna` and extra-high reasoning; all replies parsed, and no cells were redrawn.

For the sentencing study, a GPT-6 Sol run was collected on 4 October with the same 1,248 cells, prompt text, cases, displayed numbers, and displacement assignments. It used model `gpt-6-sol` at high reasoning effort. Six first attempts returned no model output and were retried once under the protocol; all final records parsed, and no tools were called. The Luna run used extra-high reasoning, so that setting differs between the two model samples.

D Design details

Crossing offense levels 16 and 24 with criminal history categories I and IV yields the four advisory ranges used throughout: 21–27, 33–41, 51–63, and 77–96 months (Figure 7).




Figure 7: (a) Anchor values at four displacements around the advisory midpoint. (b) Identical numbers under peer and forecast labels; their slope difference is the peer premium.

Write $F_c \in \{0, 1\}^4$ for case c 's factor vector and partition the sixteen design points into four groups $g_1, \dots, g_4$ subject to

$$\sum_{c \in g_k} F_{cf} = \frac{1}{2} |g_k| \quad \text{for every group } k \text{ and every factor } f, \quad (5)$$

which a backtracking search satisfies exactly at $|g_k| = 4$. Agent i then receives

$$\delta_{ic} = d_{(k(c)+i) \bmod 4}, \quad d = (-0.30, -0.15, +0.15, +0.30), \quad (6)$$

where $k(c)$ indexes the group containing c . Every agent sees each displacement exactly four times, and (5) makes the four factors split evenly within each displacement level.

E Inference

Drawing $\zeta_{ic} \sim \text{Unif}\{0, 1\}$ independently per cell and relabelling $P_{ic}^{(b)} = \zeta_{ic}P_{ic} + (1 - \zeta_{ic})(1 - P_{ic})$ generates the exact null distribution of $\hat{\Delta}$, from which

$$p = \frac{1 + \#\{b : |\hat{\Delta}^{(b)}| \geq |\hat{\Delta}|\}}{1 + B}, \quad q_{.95} = 95\text{th percentile of } |\hat{\Delta}^{(b)}|, \quad (7)$$

with $B = 9,999$.

We quote all three clusterings for the premium contrasts and the case-clustered figure for per-arm pull. Cluster-robust p -values use a wild cluster bootstrap with Webb six-point weights [7, 19, 25], which we use because agent-clustered inference rests on six clusters. Per-arm pull uses a second randomization null, permuting the displacement within each agent across the design's own factor-balanced case groups; (7) describes the attribution-label null used for premiums. Population agreement is summarised by the between-agent spread on a common case,

$$\bar{\sigma} = \frac{1}{|\mathcal{C}|} \sum_{c \in \mathcal{C}} \text{sd}_i(s_{ic}/m_c), \quad (8)$$

the mean over cases of the standard deviation across agents of the sentence expressed in units of the case midpoint.

Across 160 cells in which the same agent saw the same case under more than one attribution, the displayed numbers were identical in every paired cell. The four case factors remain balanced across displacement levels, with the largest departure from an even split equal to 0.010. Adding offense-type fixed effects leaves the guideline-present premium at -0.018 (SE 0.042).

Adding a guideline indicator G with $\delta \times G$, $P \times G$ and $\delta \times P \times G$ terms gives a moderation model whose triple interaction is 0.01 (SE 0.10, $p = 0.91$). Only arms with the original closing lines were run under both guideline conditions, and their peer-premium estimates are near zero; we therefore report the conditions separately.

F Pull by arm

Arm-level pulls include all arm agents; paired premiums use shared cells (Table 1).

Table 1: Opus 5 pull by block and source. Case-clustered SEs; randomization p -values. No-number rows show mean deviation.

Numbers attributed to	Block variant	n	Pull $\hat{\pi}$	SE	p
<i>Guideline removed</i>					
Other judges of this bench	none	64	0.88	0.06	—
Other judges of this bench	structure-matched	96	0.76	0.07	—
Other judges of this bench	paraphrase one	48	0.59	0.11	—
Other judges of this bench	paraphrase two	48	0.72	0.11	—
Statistical forecast	none	96	0.55	0.09	—
Other judges of this bench	“decide independently”	64	0.61	0.09	<0.001
Statistical forecast	“advisory only”	64	0.61	0.08	<0.001
Docketing artifact	“no legal significance”	64	0.07	0.09	0.690
<i>Guideline present</i>					
Other judges of this bench	“decide independently”	96	0.06	0.04	0.513
Statistical forecast	“advisory only”	96	0.08	0.03	0.022
No numbers shown, guideline present		48	mean deviation -0.20		
No numbers shown, guideline removed		62	mean deviation 0.26		

G Cross-model estimates

With a guideline and no numbers, between-agent spread is 0.029; without the guideline it is 0.286. Leave-one-out estimates for the no-guideline condition range from 0.074 to 0.327, so the estimated magnitude varies with the four-agent sample. Guideline-present premiums per model are -0.02 (SE 0.04 on case, $n = 192$) on Opus 5, 0.20 (SE 0.06) on Sonnet 5 and 0.21 (SE 0.11) on Haiku 4.5, the latter two over 32 decisions from one agent each.

We ran both displaced-anchor arms on three models of the same family, with and without the advisory guideline. With the guideline present, pull increases as model size decreases, from near zero on Opus 5 to roughly two-thirds on Haiku 4.5. Without the guideline, the spread in pull across the three models is smaller, although Opus 5 remains below full adoption.

The estimated peer premiums are 0.32 on Opus 5, 0.12 on Haiku 4.5, and 0.02 on Sonnet 5. The smaller-model estimates use 64 decisions each and are too imprecise to separate. We make no cross-model scaling claim from these three estimates.

The GPT-6 Luna and Sol premiums use the same cases and prompts. Their full estimates are reported in Table 2; the comparisons across models are descriptive because the reasoning settings differ.

Table 2: GPT-6 peer premiums relative to the bare forecast. Each contrast uses 416 observations per arm and case-clustered 95% intervals. The two runs used identical prompts and allocations.

Model	Effort	Contrast	Premium [95% CI]	p	MDE	n
GPT-6 Luna	xhigh	Structure-matched	0.220 [0.125, 0.316]	< 0.001	0.136	832
GPT-6 Luna	xhigh	Bare	0.256 [0.159, 0.353]	< 0.001	0.138	832
GPT-6 Sol	high	Structure-matched	0.239 [0.134, 0.344]	< 0.001	0.150	832
GPT-6 Sol	high	Bare	0.113 [0.028, 0.199]	< 0.001	0.123	832

H Confidence and exact matching

Confidence uses 4 of its ten values, and 98% of ratings are six or seven. We treat confidence as an uncalibrated self-report and describe it across arm means. Against the same bare forecast, the first independence paraphrase yields 0.05 ($p = 0.424$), while the second yields 0.18 ($p = 0.012$) (Section 4.5; Figure 8).




Figure 8: (a) Mean confidence against pull for each model-arm cell; the dashed line and correlation ($r = -0.82$) use only the six Opus 5 arm means (arm $n = 16-416$). (b) Equal-weighted mean, by provider, of model-level shares matching one of the three displayed numbers, for the three arms common to all five models.

I The cascade floors

Shuffling independent agents into pseudo-speaking orders yields a 48% exact-match rate and a slope of 0.75 ([0.57, 0.97]). Round-number clustering and case-level variation can produce these matches without peer influence. With case fixed effects, the same placebo returns -0.97 ($[-2.01, -0.13]$), consistent with the negative sign Angrist [3] predicts when case dummies place predecessor outcomes on both sides of a within-case comparison. In the cascade, the fixed-effect slope is 0.55 (SE 0.06), compared with the plain slope 1.00 (SE 0.05). The fixed-effect estimate is farther from its placebo floor; the plain estimate is close to 0.75. The placebo uses at most three predecessors, compared with five in the cascade, so it does not provide a floor matched to the later seats. The reflection problem [15] does not arise because each agent’s predecessors are fixed before it decides.

The first four agents speak in an order randomised per case; the last two were added afterwards in fixed positions, so position effects are identified over the first four. Exact reproduction rates at the second, third and fourth positions are 50%, 56% and 88%, respectively, with sixteen decisions at each position. In the docketing arm, which shows three numbers in the same layout, the exact-match rate is 8%. The final-pair spread is 0.017, compared with 0.029 for the independent population with a guideline and no peer information. The estimates use different agent sets, so we do not calculate a fold-change.

J Robustness of the premium

J.1 Design checks

Across 160 cells in which the same agent saw the same case under more than one attribution, the numbers shown differed in 0 cells, so the yoke is exact. The four case factors remain balanced across displacement levels, with the largest departure from an even split equal to 0.010. Adding offense-type fixed effects leaves the guideline-present premium at -0.018 (SE 0.042).

J.2 Sentence values and displayed numbers

Over 3807 decisions agents use 93 distinct sentence values, 65% divisible by six months, 42% by twelve and 27% by five, with 48 months the single most common answer. Of the numbers shown

to agents, 16.7% are divisible by six and 16.7% by five, and no record shows three round numbers at once. Decisions reproduce a displayed number in 29% of cases (1108 decisions), of which 45% are divisible by six; matching runs 40% where some displayed number is round against 18% where none is. The original-sentence comparison is not statistically distinguishable from zero with round numbers excluded, -0.12 ($p = 0.312$, $n = 64$). Exact matching is more common when at least one displayed number is round than when none is. We therefore compare copying with a shuffled floor. Because both attribution arms show the same numbers, round-number matching does not change the estimated premium.

J.3 Premium by displacement magnitude

We estimate the premium at each displacement. The bare premium is flat, 0.35 at $|\delta| = 0.15$ and 0.32 at $|\delta| = 0.30$; the structure-matched premium rises from 0.02 ($p = 0.869$) to 0.25 ($p = 0.005$). It is larger at the far than at the near displacement. A single slope averages these displacement-specific estimates. The original-sentence comparison is null at both magnitudes: -0.14 ($p = 0.578$) and 0.03 ($p = 0.780$).

J.4 Alternative outcome scales

On the log-ratio scale, the matched and bare premiums are 0.17 ($p = 0.029$) and 0.31 ($p < 0.001$). On the within-case midrank scale, they are 0.25 ($p = 0.040$) and 0.55 ($p < 0.001$). The original-closing comparison remains null on both scales (0.01, $p = 0.892$; -0.01 , $p = 0.982$). Since ranks and ratios use different scales, we compare signs and p -values, not magnitudes.

K Behaviour under the equivalence sentence

The dispersion and copying shifts in Section 4.8 are post-hoc diagnostics, not pre-registered. On Opus 5, outcome standard deviations are 0.235 and 0.242 in the structure-matched and bare pairs, versus 0.336 and 0.329 with the equivalence and reliability sentences. Exact displayed-number matches are 29%, 25%, and 28% in the earlier arms, versus 8%, 12%, 6%, and 8% in the four later arms. Sentences outside the span of the three numbers rise from 31–57% to 80–84%.

In both peer and forecast arms, agents independently said the numbers added no evidence beyond the case file. One wrote that they “derive from the identical case file and add no new information about the defendant, so I gave them no weight as evidence.” Responses were recorded verbatim and not re-drawn.

L Case factors

Without numbers, severity and prior record increase sentence length by 101% and 74% (Figure 9). With a guideline, 54% of decisions fall within its range, 45% below, and 1% above; departures are almost entirely downward. Proportionality is the most common stated rationale with a guideline (49%); deterrence is the most common without it (47%). The factor analysis estimates larger effects for severity and prior record (152%, 63%) than for assistance and remorse (-21% , 0%).




Figure 9: Case-factor effects on sentence length by guideline condition.

M What each arm can exclude

Of 20 estimates, 3 miss the criterion; all misses are in smaller-model cells (Table 3).

Table 3: Minimum detectable effect (MDE) is $2.80 \times \text{SE}$ for 80% power at two-sided 5%. “Powered” indicates whether each estimate meets its Claude benchmark.

Estimate	<i>n</i>	Est.	SE	MDE	Powered
Structure-matched premium, Opus 5	192	+0.206	0.079	0.222	yes
Bare-block premium, Opus 5	128	+0.323	0.047	0.132	yes
Cost of the peer closing sentence	128	-0.279	0.074	0.208	yes
Cost of paraphrase one	96	-0.326	0.066	0.184	yes
Cost of paraphrase two	96	-0.198	0.070	0.196	yes
Premium with both original closings	128	-0.006	0.068	0.191	yes
Premium, Opus 5, guideline present	192	-0.018	0.041	0.116	yes
Premium, Opus 5, guideline removed	128	-0.006	0.068	0.191	yes
Premium, Sonnet 5, guideline present	32	+0.204	0.056	0.156	yes
Premium, Sonnet 5, guideline removed	32	+0.065	0.290	0.812	no
Premium, Haiku 4.5, guideline present	32	+0.211	0.112	0.315	yes
Premium, Haiku 4.5, guideline removed	32	+0.071	0.119	0.333	no
Premium, smaller models pooled	128	+0.138	0.114	0.318	yes
GPT-6 Luna, structure-matched	832	+0.220	0.049	0.136	yes
GPT-6 Luna, bare blocks	832	+0.256	0.049	0.138	yes
GPT-6 Sol, structure-matched	832	+0.239	0.054	0.150	yes
GPT-6 Sol, bare blocks	832	+0.113	0.044	0.123	yes
Extra pull from removing guideline, Opus 5	320	+0.541	0.071	0.199	yes
Extra pull from removing guideline, Sonnet 5	64	+0.632	0.152	0.426	no
Extra pull from removing guideline, Haiku 4.5	64	+0.357	0.063	0.176	yes

N Civil-damages results across model lineages

We applied the displaced-number design to sixteen fictional civil claims, crossing severity, liability, mitigation, and documentation. D1 used Claude Opus 5. Its proximity measures were selected after examining the first eight agents, so those panel estimates are exploratory; the measures and direction were then registered before the two-agent confirmation sample (A12 and A01) was examined. On those 32 matched pairs, in-band awards were 14/32 under the peer label and 26/32 under the tool label ($p = 0.0042$, two-sided exact sign test), awards within two percent of the central figure 10/32 and 24/32 ($p = 0.00052$), and exact central matches 2/32 and 23/32 ($p = 9.54 \times 10^{-7}$). The peer-minus-tool difference in proportional distance was +0.140 (95% t [0.032, 0.248], MDE 0.142). The registered rule calls this suggestive; see `damages/AMENDMENT-D1-proximity.md`.

D2 was registered before collection and used the same ten agents, claims, primary arms, and rendered prompts as the complete D1 panel; every prompt hash matched its D1 counterpart. GPT-6 Sol at high reasoning returned 319 of 320 planned responses parseably (the A07 step-3 tool cell gave none and was not retried), leaving 159 matched pairs. Table 4 compares all ten D1 agents with D2; its D1 column pools the exploratory panel with the confirmation sample and is descriptive. Differ-

ences are peer-minus-tool with 95% claim-clustered t intervals; binary p -values are exact sign tests on discordant pairs, reported peer-only/tool-only.

Table 4: Civil-damages proximity outcomes by model run.

Outcome	Claude Opus 5, D1 (160 pairs)	GPT-6 Sol high, D2 (159 pairs)
Inside band	Peer 21/160; tool 88/160; $\Delta = -0.419$ [$-0.557, -0.281$]; $p = 2.17 \times 10^{-18}$; discordant peer-only/tool-only 2/69	Peer 90/159; tool 88/159; $\Delta = +0.013$ [$-0.127, +0.152$]; $p = 0.871$; discordant 20/18
Distance from center	Mean peer 0.392; tool 0.156; $\Delta = +0.236$ [$+0.145, +0.328$]; MDE 0.120	Mean peer 0.274; tool 0.182; $\Delta = +0.093$ [$+0.011, +0.174$]; MDE 0.107
Within 2% of center	Peer 11/160; tool 68/160; $\Delta = -0.356$ [$-0.485, -0.227$]; $p = 2.08 \times 10^{-16}$; discordant 1/58	Peer 72/159; tool 73/159; $\Delta = -0.006$ [$-0.101, +0.089$]; $p = 1.000$; discordant 17/18
Equals center	Peer 2/160; tool 49/160; $\Delta = -0.294$ [$-0.390, -0.198$]; $p = 1.42 \times 10^{-14}$; discordant 0/47	Peer 34/159; tool 43/159; $\Delta = -0.057$ [$-0.152, +0.039$]; $p = 0.200$; discordant 15/24

The registered D2 primary outcome does not reproduce D1’s large in-band difference: the contrast is small, reverses direction, and is not statistically decisive. This is not evidence that the arms are equivalent. For central distance, the D2 interval excludes zero, but its MDE exceeds the point estimate, so the registered rule calls the result suggestive rather than decisive. Neither binary secondary outcome rejects. The strong D1 proximity pattern therefore does not generalize unchanged to this GPT-6 Sol sample.

For continuity, the pooled displacement slope is $+0.383$ on D1 (95% z interval $[0.185, 0.581]$, MDE 0.283, within-agent permutation $p = 0.0014$) and $+0.559$ on D2 (95% z interval $[0.344, 0.774]$, MDE 0.307, $p = 0.0002$). The tool-minus-peer slope premium is $+0.215$ on D1 (MDE 0.311, $p = 0.0712$) and $+0.232$ on D2 (MDE 0.292, $p = 0.0341$); both estimates are below their MDEs. These secondary summaries do not replace the registered proximity outcomes. D1 used Claude subagents, whereas D2 used Codex CLI, so the cross-run difference is descriptive across both model and execution route. The results do not separate trust in a source label from adoption of a supplied number.